

Atlas Reflection

Key Learnings

The main lesson from Atlas was that the quality of an AI application is decided before the first line of code. AI coding tools can scaffold pages, routes, tests, and documentation quickly. That makes implementation cheaper, but it also makes it possible to build the wrong product faster. I found that the most valuable work was early: defining a real problem, choosing a narrow user, comparing alternatives, and writing down what the product should and should not do.

I began with validation instead of features. The seven-step framework helped me turn a broad career-coach idea into a specific workflow: a job seeker has a resume and target role but does not know what evidence matters or what action to take next. Research on early-career employment, changing skills, existing resume tools, and paid alternatives clarified the gap. Atlas should not try to replace job boards, ATS scanners, or mentors. Its role is to turn a role-specific review into a practical, trackable action plan.

The PRD and specification made this decision executable. The PRD set the target user, MVP, success metrics, and boundaries. The specification then defined the inputs, outputs, behavior rules, privacy constraints, architecture, and acceptance criteria. These documents made AI-assisted development more precise: instead of asking an agent to make a generic career app, I could ask it to build a small, testable part of an agreed workflow. Careful brainstorming, planning, and specifications are now a better use of time than trying to write every line manually. Code is cheaper to produce; product judgment is still scarce.

Challenges and Scope Control

The largest challenge was controlling scope within the capstone timeline. It was easy to imagine LinkedIn integration, job-board search, automatic applications, broad career courses, social sharing, and a full game system. Each idea sounded useful, but together they would have prevented a reliable MVP. I kept Atlas centered on one path: sign in, add a resume and target role, review the extracted text, generate a readiness report, complete a quest, and ask a report-specific follow-up question.

Technical scope control mattered as much as product scope. The implementation combined document extraction, structured LLM output, curated RAG guidance, Supabase Auth and Row Level Security, a readiness dashboard, and a deployed document service. I treated privacy and reliability as product requirements: uploaded files and full raw resume text are not stored, private career data is not added to the shared RAG corpus, and the fit score is guidance rather than a hiring prediction. Testing validation, schemas, permissions, document safety, and deployment paths showed me that an AI feature needs more than a model call.

Future Versions

The next product priority is primary user validation: run the five-user concierge test, capture anonymized feedback, and check whether users can identify and complete a useful next action. I would also add user-controlled report deletion, a wider library of career-path presets, richer report-quality evaluation, and progress history across multiple reports.

I would make the experience more gamified in a purposeful way, not simply add points. A future Atlas could include a visual campaign map with milestone checkpoints, optional mission categories for resume evidence, portfolio proof, networking, and interview practice, and a “next best mission” based on the user’s current gap. Completing a mission could unlock a private progress recap, a portfolio evidence checklist, or a new campaign stage. XP, ranks, badges, and level-up moments should remain private and tied to real evidence created by the user. I would avoid public leaderboards, streak loss, countdown pressure, or rewards that imply completing tasks guarantees an interview or job offer.

Overall, Atlas showed me how validation, planning, implementation, testing, and deployment fit together. AI can accelerate delivery, but I remain responsible for deciding what is worth building, keeping the scope realistic, and checking that the final product is safe and useful.